

Data Structures

Atli FF

15. september 2023

School of Computer Science

Reykjavík University

Today's material

- Built-in data structures and their applications
- Augmenting a data structure
- Union-Find
- Precomputations like prefix sums
- Square root decomposition
- Segment trees
- Sparse tables

Applications of Arrays and Linked Lists

- Too many to list, more so for arrays

Applications of Arrays and Linked Lists

- Too many to list, more so for arrays
- Most problems require storing data, usually in an array

Applications of Arrays and Linked Lists

- Too many to list, more so for arrays
- Most problems require storing data, usually in an array
- On modern machines, arrays are almost always a better choice than a linked list

Applications of Arrays and Linked Lists

- Too many to list, more so for arrays
- Most problems require storing data, usually in an array
- On modern machines, arrays are almost always a better choice than a linked list
- There are however a few cases where linked lists are better

Example problem: Broken Keyboard

- <http://uva.onlinejudge.org/external/119/11988.html>

- Processing events in a last-in first-out order

Applications of Stacks

- Processing events in a last-in first-out order
- Simulating recursion

Applications of Stacks

- Processing events in a last-in first-out order
- Simulating recursion
- Depth-first search in a graph

Applications of Stacks

- Processing events in a last-in first-out order
- Simulating recursion
- Depth-first search in a graph
- Reverse a sequence

Applications of Stacks

- Processing events in a last-in first-out order
- Simulating recursion
- Depth-first search in a graph
- Reverse a sequence
- Matching brackets

Applications of Stacks

- Processing events in a last-in first-out order
- Simulating recursion
- Depth-first search in a graph
- Reverse a sequence
- Matching brackets
- And a lot more

Applications of Queues

- Processing events in a first-in first-out order

Applications of Queues

- Processing events in a first-in first-out order
- Breadth-first search in a graph

Applications of Queues

- Processing events in a first-in first-out order
- Breadth-first search in a graph
- And a lot more

Applications of Priority Queues

- Processing events in order of priority

Applications of Priority Queues

- Processing events in order of priority
- Finding a shortest path in a graph

Applications of Priority Queues

- Processing events in order of priority
- Finding a shortest path in a graph
- Some greedy algorithms

Applications of Priority Queues

- Processing events in order of priority
- Finding a shortest path in a graph
- Some greedy algorithms
- And a lot more

Applications of Sets

- Keep track of distinct items

Applications of Sets

- Keep track of distinct items
- Have we seen an item before?

Applications of Sets

- Keep track of distinct items
- Have we seen an item before?
- If implemented as a binary search tree:

Applications of Sets

- Keep track of distinct items
- Have we seen an item before?
- If implemented as a binary search tree:
 - Find the successor/predecessor of an element

Applications of Sets

- Keep track of distinct items
- Have we seen an item before?
- If implemented as a binary search tree:
 - Find the successor/predecessor of an element
 - Count how many elements are less than a given element

Applications of Sets

- Keep track of distinct items
- Have we seen an item before?
- If implemented as a binary search tree:
 - Find the successor/predecessor of an element
 - Count how many elements are less than a given element
 - Count how many elements are between two given elements

Applications of Sets

- Keep track of distinct items
- Have we seen an item before?
- If implemented as a binary search tree:
 - Find the successor/predecessor of an element
 - Count how many elements are less than a given element
 - Count how many elements are between two given elements
 - Find the k th largest element

Applications of Sets

- Keep track of distinct items
- Have we seen an item before?
- If implemented as a binary search tree:
 - Find the successor/predecessor of an element
 - Count how many elements are less than a given element
 - Count how many elements are between two given elements
 - Find the k th largest element
- And a lot more

Applications of Maps

- Associating a value with a key

Applications of Maps

- Associating a value with a key
- As a frequency table

Applications of Maps

- Associating a value with a key
- As a frequency table
- As a memory when doing Dynamic Programming

Applications of Maps

- Associating a value with a key
- As a frequency table
- As a memory when doing Dynamic Programming
- And a lot more

Augmenting Data Structures

- Sometimes we can store extra information in our data structures to gain more functionality

Augmenting Data Structures

- Sometimes we can store extra information in our data structures to gain more functionality
- Usually we can't do this to data structures in the standard library

Augmenting Data Structures


- Sometimes we can store extra information in our data structures to gain more functionality
- Usually we can't do this to data structures in the standard library
- Need our own implementation that we can customize

Augmenting Data Structures

- Sometimes we can store extra information in our data structures to gain more functionality
- Usually we can't do this to data structures in the standard library
- Need our own implementation that we can customize
- Example: Augmenting binary search trees


Augmenting Binary Search Trees

- We have a binary search tree and want to efficiently:
 - Count number of elements $< x$
 - Find the k th smallest element
- Naive method is to go through all vertices, but that is slow: $O(n)$




Augmenting Binary Search Trees

- Idea: In each vertex store the size of the subtree
- This information can be maintained when we insert/delete elements without increasing time complexity




Augmenting Binary Search Trees

- Count number of elements < 38
 - Search for 38 in the tree
 - Count the vertices that we pass by that are less than x
 - When we are at a vertex where we should go right, get the size of the left subtree and add it to our count




Augmenting Binary Search Trees

- Count number of elements < 38
 - Search for 38 in the tree
 - Count the vertices that we pass by that are less than x
 - When we are at a vertex where we should go right, get the size of the left subtree and add it to our count
- Time complexity $O(\log n)$




Augmenting Binary Search Trees

- Find k th smallest element
 - We're on a vertex whose left subtree is of size m
 - If $k = m + 1$, we found it
 - If $k \leq m$, look for the k th smallest element in the left subtree
 - If $k > m + 1$, look for the $k - m - 1$ st smallest element in the right subtree



Augmenting Binary Search Trees

- Find k th smallest element
 - We're on a vertex whose left subtree is of size m
 - If $k = m + 1$, we found it
 - If $k \leq m$, look for the k th smallest element in the left subtree
 - If $k > m + 1$, look for the $k - m - 1$ st smallest element in the right subtree
- Example: $k = 11$



Union-Find

- We have n items

Union-Find

- We have n items
- Maintains a collection of disjoint sets (or equivalently, an equivalence relation)

Union-Find

- We have n items
- Maintains a collection of disjoint sets (or equivalently, an equivalence relation)
- Each of the n items is in exactly one set

Union-Find

- We have n items
- Maintains a collection of disjoint sets (or equivalently, an equivalence relation)
- Each of the n items is in exactly one set
- We represent each set with one of its members, a representative element

Union-Find

- We have n items
- Maintains a collection of disjoint sets (or equivalently, an equivalence relation)
- Each of the n items is in exactly one set
- We represent each set with one of its members, a representative element
- Supports two operations efficiently: `find(x)` and `union(x,y)`.

Union-Find

- We have n items
- Maintains a collection of disjoint sets (or equivalently, an equivalence relation)
- Each of the n items is in exactly one set
- We represent each set with one of its members, a representative element
- Supports two operations efficiently: `find(x)` and `union(x,y)`.
- Operation `find(x)` finds the representative of the set x is in

Union-Find

- We have n items
- Maintains a collection of disjoint sets (or equivalently, an equivalence relation)
- Each of the n items is in exactly one set
- We represent each set with one of its members, a representative element
- Supports two operations efficiently: `find(x)` and `union(x,y)`.
- Operation `find(x)` finds the representative of the set x is in
- Operation `union(x, y)` unions the sets of which x and y are members.

Union-Find

- It is generally initialized with all items being in their own set.

Union-Find

- It is generally initialized with all items being in their own set.
- So for $n = 5$ we start out with $\{\{1\}, \{2\}, \{3\}, \{4\}, \{5\}\}$.

Union-Find

- It is generally initialized with all items being in their own set.
- So for $n = 5$ we start out with $\{\{1\}, \{2\}, \{3\}, \{4\}, \{5\}\}$.
- `join(1, 3)` then changes this to $\{\{1, 3\}, \{2\}, \{4\}, \{5\}\}$.

Union-Find

- It is generally initialized with all items being in their own set.
- So for $n = 5$ we start out with $\{\{1\}, \{2\}, \{3\}, \{4\}, \{5\}\}$.
- `join(1, 3)` then changes this to $\{\{1, 3\}, \{2\}, \{4\}, \{5\}\}$.
- `join(2, 5)` then results in $\{\{1, 3\}, \{2, 5\}, \{4\}\}$.

Union-Find

- It is generally initialized with all items being in their own set.
- So for $n = 5$ we start out with $\{\{1\}, \{2\}, \{3\}, \{4\}, \{5\}\}$.
- `join(1, 3)` then changes this to $\{\{1, 3\}, \{2\}, \{4\}, \{5\}\}$.
- `join(2, 5)` then results in $\{\{1, 3\}, \{2, 5\}, \{4\}\}$.
- `join(2, 4)` then results in $\{\{1, 3\}, \{2, 4, 5\}\}$.

Union-Find

- It is generally initialized with all items being in their own set.
- So for $n = 5$ we start out with $\{\{1\}, \{2\}, \{3\}, \{4\}, \{5\}\}$.
- $\text{join}(1, 3)$ then changes this to $\{\{1, 3\}, \{2\}, \{4\}, \{5\}\}$.
- $\text{join}(2, 5)$ then results in $\{\{1, 3\}, \{2, 5\}, \{4\}\}$.
- $\text{join}(2, 4)$ then results in $\{\{1, 3\}, \{2, 4, 5\}\}$.
- $\text{join}(1, 4)$ finally results in $\{\{1, 2, 3, 4, 5\}\}$.

Union-Find

- It is generally initialized with all items being in their own set.
- So for $n = 5$ we start out with $\{\{1\}, \{2\}, \{3\}, \{4\}, \{5\}\}$.
- `join(1, 3)` then changes this to $\{\{1, 3\}, \{2\}, \{4\}, \{5\}\}$.
- `join(2, 5)` then results in $\{\{1, 3\}, \{2, 5\}, \{4\}\}$.
- `join(2, 4)` then results in $\{\{1, 3\}, \{2, 4, 5\}\}$.
- `join(1, 4)` finally results in $\{\{1, 2, 3, 4, 5\}\}$.
- At any given point `find(x)` returns some value in the same set as x .

Union-Find

- It is generally initialized with all items being in their own set.
- So for $n = 5$ we start out with $\{\{1\}, \{2\}, \{3\}, \{4\}, \{5\}\}$.
- $\text{join}(1, 3)$ then changes this to $\{\{1, 3\}, \{2\}, \{4\}, \{5\}\}$.
- $\text{join}(2, 5)$ then results in $\{\{1, 3\}, \{2, 5\}, \{4\}\}$.
- $\text{join}(2, 4)$ then results in $\{\{1, 3\}, \{2, 4, 5\}\}$.
- $\text{join}(1, 4)$ finally results in $\{\{1, 2, 3, 4, 5\}\}$.
- At any given point $\text{find}(x)$ returns some value in the same set as x .
- The important bit is that $\text{find}(x)$ returns the same value for all elements of the same set, the representative.

Union-Find

- We can do this by maintaining an array of parents, letting the i -th value be the index of the parent of the i -th item.

Union-Find

- We can do this by maintaining an array of parents, letting the i -th value be the index of the parent of the i -th item.
- If a value has no parent, we can denote this somehow, make it its own parent, give it the value -1 , exactly what we do is not important.

Union-Find

- We can do this by maintaining an array of parents, letting the i -th value be the index of the parent of the i -th item.
- If a value has no parent, we can denote this somehow, make it its own parent, give it the value -1 , exactly what we do is not important.
- To get the representative of x we go to the parent of our current item (starting at x) until the item has no parent.

Union-Find

- We can do this by maintaining an array of parents, letting the i -th value be the index of the parent of the i -th item.
- If a value has no parent, we can denote this somehow, make it its own parent, give it the value -1 , exactly what we do is not important.
- To get the representative of x we go to the parent of our current item (starting at x) until the item has no parent.
- Then to unite x, y we simply make the representative of x the parent of the representative of y .

Naïve Union-Find implementation

```
struct union_find {  
    vector<int> parent;  
    union_find(int n) {  
        parent = vector<int>(n);  
        for(int i = 0; i < n; i++) {  
            parent[i] = i;  
        }  
    }  
    int find(int x) {  
        return parent[x] == x ? x : find(parent[x]);  
    }  
    void unite(int x, int y) {  
        parent[find(x)] = find(y);  
    }  
};
```

Union-Find

- The problem here is that for specific queries, the parent chains may end up being of length $\mathcal{O}(n)$, making each query linear.

Union-Find

- The problem here is that for specific queries, the parent chains may end up being of length $\mathcal{O}(n)$, making each query linear.
- The key to making this more efficient is making those chains shorter.

Union-Find

- The problem here is that for specific queries, the parent chains may end up being of length $\mathcal{O}(n)$, making each query linear.
- The key to making this more efficient is making those chains shorter.
- We do this by flattening the chain each time we query `find`, so the amortized complexity becomes good.

Union-Find

- The problem here is that for specific queries, the parent chains may end up being of length $\mathcal{O}(n)$, making each query linear.
- The key to making this more efficient is making those chains shorter.
- We do this by flattening the chain each time we query `find`, so the amortized complexity becomes good.
- The worst case is still $\mathcal{O}(n)$ but the amortized complexity is $\mathcal{O}(\alpha(n))$ which may as well be a constant, as it is < 5 for n equal to the number of atoms in the observable universe.

Path compressed Union-Find implementation

```
struct union_find {  
    vector<int> parent;  
    union_find(int n) {  
        parent = vector<int>(n);  
        for (int i = 0; i < n; i++) {  
            parent[i] = i;  
        }  
    }  
    int find(int x) {  
        if(parent[x] == x) return x;  
        return parent[x] = find(parent[x]);  
    }  
    void unite(int x, int y) {  
        parent[find(x)] = find(y);  
    }  
};
```

Union-Find applications

- Union-Find maintains a collection of disjoint sets

Union-Find applications

- Union-Find maintains a collection of disjoint sets
- When are we dealing with such collections?

Union-Find applications

- Union-Find maintains a collection of disjoint sets
- When are we dealing with such collections?
- Usually when we want to work with equivalence relations like graph connectivity

Union-Find applications

- Union-Find maintains a collection of disjoint sets
- When are we dealing with such collections?
- Usually when we want to work with equivalence relations like graph connectivity
- By modifying the data structure it can also contain more queryable data

Union-Find applications

- Union-Find maintains a collection of disjoint sets
- When are we dealing with such collections?
- Usually when we want to work with equivalence relations like graph connectivity
- By modifying the data structure it can also contain more queryable data
 - Number of different sets currently

Union-Find applications

- Union-Find maintains a collection of disjoint sets
- When are we dealing with such collections?
- Usually when we want to work with equivalence relations like graph connectivity
- By modifying the data structure it can also contain more queryable data
 - Number of different sets currently
 - Current size of the set containing x

Union-Find applications

- Union-Find maintains a collection of disjoint sets
- When are we dealing with such collections?
- Usually when we want to work with equivalence relations like graph connectivity
- By modifying the data structure it can also contain more queryable data
 - Number of different sets currently
 - Current size of the set containing x
 - An iterable list of all elements of the set containing x

Union-Find applications

- Union-Find maintains a collection of disjoint sets
- When are we dealing with such collections?
- Usually when we want to work with equivalence relations like graph connectivity
- By modifying the data structure it can also contain more queryable data
 - Number of different sets currently
 - Current size of the set containing x
 - An iterable list of all elements of the set containing x
- When tracking size you can use it to always perform small-to-large merges for $\mathcal{O}(\log n)$ time complexity.

Example problem: Skolavslutningen

- <https://open.kattis.com/problems/skolavslutningen>

Range queries

- We have an array A of size n .

Range queries

- We have an array A of size n .
- Given i, j , we want to answer:

Range queries

- We have an array A of size n .
- Given i, j , we want to answer:
 - $\max(A[i], A[i + 1], \dots, A[j - 1], A[j])$

Range queries

- We have an array A of size n .
- Given i, j , we want to answer:
 - $\max(A[i], A[i + 1], \dots, A[j - 1], A[j])$
 - $\min(A[i], A[i + 1], \dots, A[j - 1], A[j])$

Range queries

- We have an array A of size n .
- Given i, j , we want to answer:
 - $\max(A[i], A[i + 1], \dots, A[j - 1], A[j])$
 - $\min(A[i], A[i + 1], \dots, A[j - 1], A[j])$
 - $\text{sum}(A[i], A[i + 1], \dots, A[j - 1], A[j])$

Range queries

- We have an array A of size n .
- Given i, j , we want to answer:
 - $\max(A[i], A[i + 1], \dots, A[j - 1], A[j])$
 - $\min(A[i], A[i + 1], \dots, A[j - 1], A[j])$
 - $\text{sum}(A[i], A[i + 1], \dots, A[j - 1], A[j])$
- We want to answer these queries efficiently, or in other words, without looking through all elements.

Range queries

- We have an array A of size n .
- Given i, j , we want to answer:
 - $\max(A[i], A[i + 1], \dots, A[j - 1], A[j])$
 - $\min(A[i], A[i + 1], \dots, A[j - 1], A[j])$
 - $\text{sum}(A[i], A[i + 1], \dots, A[j - 1], A[j])$
- We want to answer these queries efficiently, or in other words, without looking through all elements.
- Sometimes we also want to update elements.

Range sum on a static array

- Let's look at range sums on a constant array

Range sum on a static array

- Let's look at range sums on a constant array
- How do we support these queries efficiently?

Range sum on a static array

- Let's look at range sums on a constant array
- How do we support these queries efficiently?
- Simplification: only support queries of the form $\text{sum}(0, j)$

Range sum on a static array

- Let's look at range sums on a constant array
- How do we support these queries efficiently?
- Simplification: only support queries of the form $\text{sum}(0, j)$
- Notice that $\text{sum}(i, j) = \text{sum}(0, j) - \text{sum}(0, i - 1)$

Range sum on a static array

- So we're only interested in prefix sums

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1						

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1					

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8				

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16			

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16	21		

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16	21	30	

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16	21	30	33

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16	21	30	33

- $O(n)$ time to preprocess

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16	21	30	33

- $O(n)$ time to preprocess
- $O(1)$ time each query

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16	21	30	33

- $O(n)$ time to preprocess
- $O(1)$ time each query
- Can we support updating efficiently?

Range sum on a static array

- So we're only interested in prefix sums
- But there are only n of them...
- Just compute them all once in the beginning

1	0	7	8	5	9	3
1	1	8	16	21	30	33

- $O(n)$ time to preprocess
- $O(1)$ time each query
- Can we support updating efficiently? No, at least not without modification

Generalizing

- This works on any invertible function.

Generalizing

- This works on any invertible function.
- If we want the product we can store the products and use $\text{mul}(i, j) = \text{mul}(0, j) / \text{mul}(0, i - 1)$.

Generalizing

- This works on any invertible function.
- If we want the product we can store the products and use $\text{mul}(i, j) = \text{mul}(0, j) / \text{mul}(0, i - 1)$.
- This also works for multidimensional arrays, but the math is more involved.

Generalizing


- This works on any invertible function.
- If we want the product we can store the products and use $\text{mul}(i, j) = \text{mul}(0, j) / \text{mul}(0, i - 1)$.
- This also works for multidimensional arrays, but the math is more involved.
- We let $\text{sum}(x_i, x_j, y_i, y_j)$ denote the query for the sum from x_i to x_j along the x -dimension, and the same for y .

Generalizing

- This works on any invertible function.
- If we want the product we can store the products and use $\text{mul}(i, j) = \text{mul}(0, j) / \text{mul}(0, i - 1)$.
- This also works for multidimensional arrays, but the math is more involved.
- We let $\text{sum}(x_i, x_j, y_i, y_j)$ denote the query for the sum from x_i to x_j along the x -dimension, and the same for y .
- Then the formula becomes


$$\begin{aligned}\text{sum}(x_i, x_j, y_i, y_j) &= \text{sum}(0, x_j, 0, y_j) \\ &\quad - \text{sum}(0, x_{i-1}, 0, y_j) \\ &\quad - \text{sum}(0, x_j, 0, y_{i-1}) \\ &\quad + \text{sum}(0, x_{i-1}, 0, y_{i-1})\end{aligned}$$

2D sum



`query(1, 3, 3, 4)`


2D sum



`query(1, 3, 3, 4)`

`query(0, 3, 0, 4)`

2D sum




query(1, 3, 3, 4)

query(0, 3, 0, 4)

query(0, 4, 0, 2)

2D sum




query(1, 3, 3, 4)

query(0, 3, 0, 4)

query(0, 4, 0, 2)

query(0, 0, 0, 4)

2D sum



query(1, 3, 3, 4)

query(0, 3, 0, 4)

query(0, 4, 0, 2)

query(0, 0, 0, 4)

query(0, 0, 0, 2)

Range sum on a mutable array

- What if we want to support:

Range sum on a mutable array

- What if we want to support:
 - sum over a range

Range sum on a mutable array

- What if we want to support:
 - sum over a range
 - updating an element

Range sum on a mutable array

- What if we want to support:
 - sum over a range
 - updating an element
- How do we support these queries efficiently?

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket
- Time complexity: $O(k)$

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket
- Time complexity: $O(k)$
- Again we want to query over a range

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket
- Time complexity: $O(k)$
- Again we want to query over a range
 - When a bucket is contained in the range, use the stored sum for the bucket

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket
- Time complexity: $O(k)$
- Again we want to query over a range
 - When a bucket is contained in the range, use the stored sum for the bucket
 - This (sometimes) allows us to “jump” over intervals of size k

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket
- Time complexity: $O(k)$
- Again we want to query over a range
 - When a bucket is contained in the range, use the stored sum for the bucket
 - This (sometimes) allows us to “jump” over intervals of size k
 - Only have to go inside at most two buckets (each end)

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket
- Time complexity: $O(k)$
- Again we want to query over a range
 - When a bucket is contained in the range, use the stored sum for the bucket
 - This (sometimes) allows us to “jump” over intervals of size k
 - Only have to go inside at most two buckets (each end)
 - Have to consider at most n/k buckets and 2 buckets of size k

First attempt: Buckets

- Group values into buckets of size k and store result of each bucket
- Updating is easy:
 - change the array element
 - recompute corresponding bucket
- Time complexity: $O(k)$
- Again we want to query over a range
 - When a bucket is contained in the range, use the stored sum for the bucket
 - This (sometimes) allows us to “jump” over intervals of size k
 - Only have to go inside at most two buckets (each end)
 - Have to consider at most n/k buckets and 2 buckets of size k
- Time complexity: $O(n/k + k)$

Buckets: Choosing k

- Now we have a data structure that supports:
 - Querying in $O(n/k + k)$

Buckets: Choosing k

- Now we have a data structure that supports:
 - Updating in $O(k)$
 - Querying in $O(n/k + k)$

Buckets: Choosing k

- Now we have a data structure that supports:
 - Updating in $O(k)$
 - Querying in $O(n/k + k)$
- What k to pick?

Buckets: Choosing k

- Now we have a data structure that supports:
 - Updating in $O(k)$
 - Querying in $O(n/k + k)$
- What k to pick?
- Time complexity is minimized for $k = \sqrt{n}$:

Buckets: Choosing k

- Now we have a data structure that supports:
 - Updating in $O(k)$
 - Querying in $O(n/k + k)$
- What k to pick?
- Time complexity is minimized for $k = \sqrt{n}$:
 - Updating in $O(\sqrt{n})$

Buckets: Choosing k

- Now we have a data structure that supports:
 - Updating in $O(k)$
 - Querying in $O(n/k + k)$
- What k to pick?
- Time complexity is minimized for $k = \sqrt{n}$:
 - Updating in $O(\sqrt{n})$
 - Querying in $O(n/\sqrt{n} + \sqrt{n}) = O(\sqrt{n})$

Buckets: Choosing k

- Now we have a data structure that supports:
 - Updating in $O(k)$
 - Querying in $O(n/k + k)$
- What k to pick?
- Time complexity is minimized for $k = \sqrt{n}$:
 - Updating in $O(\sqrt{n})$
 - Querying in $O(n/\sqrt{n} + \sqrt{n}) = O(\sqrt{n})$
- Also known as square root decomposition, and is a very powerful technique

Example problem: Supercomputer

- <https://open.kattis.com/problems/supercomputer>

Range queries

- Now we know how to do these queries in $O(\sqrt{n})$

Range queries

- Now we know how to do these queries in $O(\sqrt{n})$
- May be too slow if n is large and many queries

Range queries

- Now we know how to do these queries in $O(\sqrt{n})$
- May be too slow if n is large and many queries
- Can we do better?

Second attempt: Segment Tree

- We create a perfect binary tree where the leaves are the elements of the array.

Second attempt: Segment Tree

- We create a perfect binary tree where the leaves are the elements of the array.
- Then each internal vertex is the sum of the values below it.

Second attempt: Segment Tree

- We create a perfect binary tree where the leaves are the elements of the array.
- Then each internal vertex is the sum of the values below it.
- Then we have $\mathcal{O}(n)$ nodes and each query can be pieced together from $\mathcal{O}(\log(n))$ node values.


Second attempt: Segment Tree

- We create a perfect binary tree where the leaves are the elements of the array.
- Then each internal vertex is the sum of the values below it.
- Then we have $\mathcal{O}(n)$ nodes and each query can be pieced together from $\mathcal{O}(\log(n))$ node values.
- We travel down the tree looking for the left and right end points, adding intervals that are completely inside our query range.


Second attempt: Segment Tree

- We create a perfect binary tree where the leaves are the elements of the array.
- Then each internal vertex is the sum of the values below it.
- Then we have $\mathcal{O}(n)$ nodes and each query can be pieced together from $\mathcal{O}(\log(n))$ node values.
- We travel down the tree looking for the left and right endpoints, adding intervals that are completely inside our query range.
- When we update a value we only need to update the parents of that node up to the root, at most $\mathcal{O}(\log(n))$ nodes.

Drawn Segment Tree, $n = 4$



Drawn Segment Tree, $n = 7$




Segment Tree - Code


```
struct segment_tree {
    segment_tree *left, *right;
    int from, to, value;
    segment_tree(int from, int to)
        : from(from), to(to), left(NULL), right(NULL), value(0) { }
};
```

```
segment_tree* build(const vector<int> &arr, int l, int r) {
    if (l > r) return NULL;
    segment_tree *res = new segment_tree(l, r);
    if (l == r) {
        res->value = arr[l];
    } else {
        int m = (l + r) / 2;
        res->left = build(arr, l, m);
        res->right = build(arr, m + 1, r);
        if (res->left != NULL) res->value += res->left->value;
        if (res->right != NULL) res->value += res->right->value;
    }
    return res;
}
```


Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates



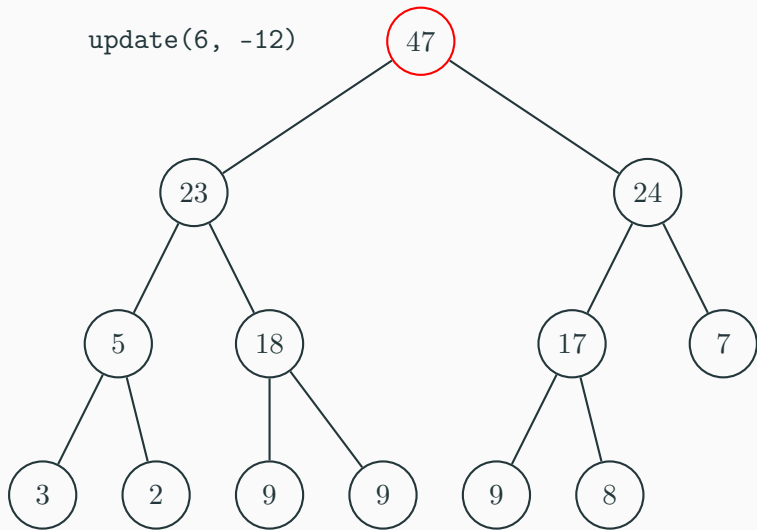
Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates




Updates



Updates




Updates




Updating a Segment Tree - Code

```
int update(segment_tree *tree, int i, int val) {
    if (tree == NULL) return 0;
    if (tree->to < i) return tree->value;
    if (i < tree->from) return tree->value;
    if (tree->from == tree->to && tree->from == i) {
        tree->value = val;
    } else {
        tree->value = update(tree->left, i, val) + update(tree->right, i, val);
    }
    return tree->value;
}
```


Querying




Querying




Querying




Querying




Querying



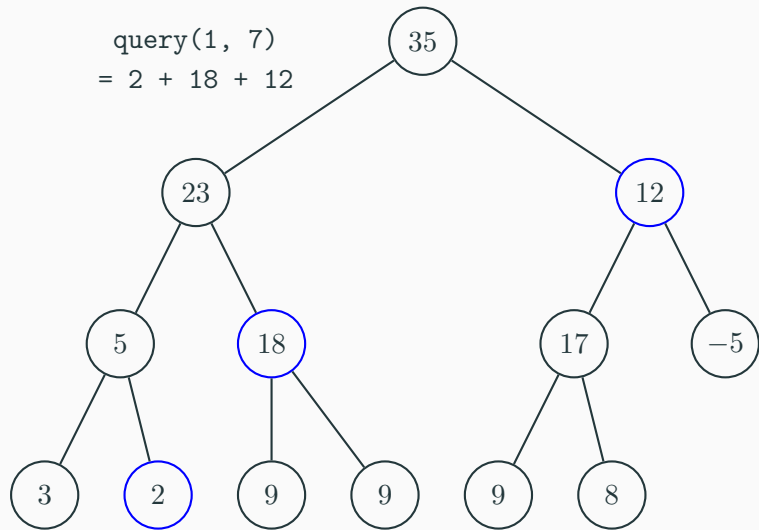
Querying



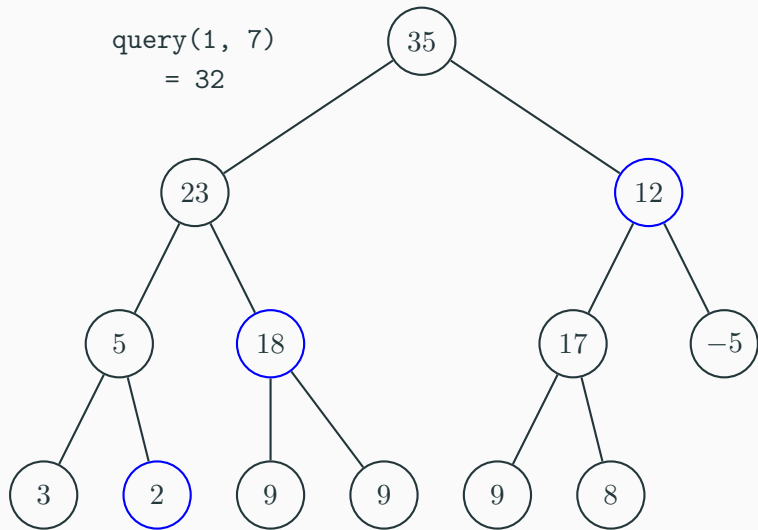
Querying




Querying




Querying




Querying




Querying




Querying




Querying




Querying




Querying



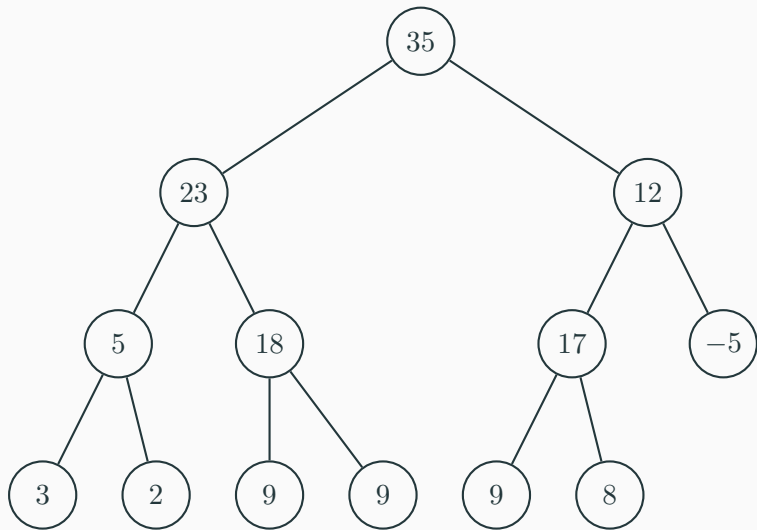
Querying



Querying



Querying



Querying a Segment Tree - Code

```
int query(segment_tree *tree, int l, int r) {  
    if (tree == NULL) return 0;  
    if (l <= tree->from && tree->to <= r) return tree->value;  
    if (tree->to < l) return 0;  
    if (r < tree->from) return 0;  
    return query(tree->left, l, r) + query(tree->right, l, r);  
}
```

Segment Tree

- Simple to use Segment Trees for min, max, gcd, and other similar operators, basically the same code.

Segment Tree

- Simple to use Segment Trees for min, max, gcd, and other similar operators, basically the same code.
- Any associative operator will work.

Segment Tree

- Simple to use Segment Trees for min, max, gcd, and other similar operators, basically the same code.
- Any associative operator will work.
- So any operator f such that $f(a, f(b, c)) = f(f(a, b), c)$ for all a, b, c .

Segment Tree

- Simple to use Segment Trees for min, max, gcd, and other similar operators, basically the same code.
- Any associative operator will work.
- So any operator f such that $f(a, f(b, c)) = f(f(a, b), c)$ for all a, b, c .
- Also possible to update a range of values in $O(\log n)$, which will be covered in bonus slides.

Example problem: Movie Collection

- <https://open.kattis.com/problems/moviecollection>

Another $\log(n)$ idea

- What if we tried something more akin to an array.

Another $\log(n)$ idea

- What if we tried something more akin to an array.
- Could we store $\log(n)$ amounts of data per element somehow?

Another $\log(n)$ idea

- What if we tried something more akin to an array.
- Could we store $\log(n)$ amounts of data per element somehow?
- Yes! For each i we can store the sum on the interval $[i, i + 2^j - 1]$ for \log many j .

Another $\log(n)$ idea

- What if we tried something more akin to an array.
- Could we store $\log(n)$ amounts of data per element somehow?
- Yes! For each i we can store the sum on the interval $[i, i + 2^j - 1]$ for \log many j .
- Then to retrieve a sum from i to j we always take the biggest chunk we can that's stored at i , which will always be at least half.

Another $\log(n)$ idea

- What if we tried something more akin to an array.
- Could we store $\log(n)$ amounts of data per element somehow?
- Yes! For each i we can store the sum on the interval $[i, i + 2^j - 1]$ for \log many j .
- Then to retrieve a sum from i to j we always take the biggest chunk we can that's stored at i , which will always be at least half.
- Then we continue until we reach j , moving i along and collecting the results.

Another $\log(n)$ idea

- What if we tried something more akin to an array.
- Could we store $\log(n)$ amounts of data per element somehow?
- Yes! For each i we can store the sum on the interval $[i, i + 2^j - 1]$ for \log many j .
- Then to retrieve a sum from i to j we always take the biggest chunk we can that's stored at i , which will always be at least half.
- Then we continue until we reach j , moving i along and collecting the results.
- This is what is known as a sparse table.

Sparse tables

- Calculating all of these values takes $\mathcal{O}(n \log(n))$ because we can calculate the values in order of increasing j .

Sparse tables

- Calculating all of these values takes $\mathcal{O}(n \log(n))$ because we can calculate the values in order of increasing j .
- Then when we calculate the sum of $[i, i + 2^j - 1]$ we just combine the earlier results of $[i, i + 2^{j-1} - 1]$ and $[i + 2^{j-1}, i + 2^j - 1]$.

Sparse tables

- Calculating all of these values takes $\mathcal{O}(n \log(n))$ because we can calculate the values in order of increasing j .
- Then when we calculate the sum of $[i, i + 2^j - 1]$ we just combine the earlier results of $[i, i + 2^{j-1} - 1]$ and $[i + 2^{j-1}, i + 2^j - 1]$.
- Querying takes $\mathcal{O}(\log(n))$, however updating is slow and difficult.

Sparse tables

- Calculating all of these values takes $\mathcal{O}(n \log(n))$ because we can calculate the values in order of increasing j .
- Then when we calculate the sum of $[i, i + 2^j - 1]$ we just combine the earlier results of $[i, i + 2^{j-1} - 1]$ and $[i + 2^{j-1}, i + 2^j - 1]$.
- Querying takes $\mathcal{O}(\log(n))$, however updating is slow and difficult.
- Why would we then ever use this instead of segment trees?

Binary lifting

- The reason might be is that with sparse tables we can do many things that segment trees can not because of how the results are combined.

Binary lifting

- The reason might be is that with sparse tables we can do many things that segment trees can not because of how the results are combined.
- Let us consider binary lifting in particular.

Binary lifting

- The reason might be is that with sparse tables we can do many things that segment trees can not because of how the results are combined.
- Let us consider binary lifting in particular.
- Suppose we have some function f that rearranges the values $\{0, 1, \dots, n - 1\}$ and we get q queries asking what happens to x if we apply f exactly m times to x .

Binary lifting

- The reason might be is that with sparse tables we can do many things that segment trees can not because of how the results are combined.
- Let us consider binary lifting in particular.
- Suppose we have some function f that rearranges the values $\{0, 1, \dots, n - 1\}$ and we get q queries asking what happens to x if we apply f exactly m times to x .
- The naïve solution is to calculate it every time, giving a time complexity of $\mathcal{O}(qm\mathcal{O}(f))$.

Binary lifting

- The reason might be is that with sparse tables we can do many things that segment trees can not because of how the results are combined.
- Let us consider binary lifting in particular.
- Suppose we have some function f that rearranges the values $\{0, 1, \dots, n - 1\}$ and we get q queries asking what happens to x if we apply f exactly m times to x .
- The naïve solution is to calculate it every time, giving a time complexity of $\mathcal{O}(qm\mathcal{O}(f))$.
- How might we use sparse tables to do better?

Binary lifting ctd.

- Let $f^{[y]}(x)$ denote the result of applying f exactly y times to x

Binary lifting ctd.

- Let $f^{[y]}(x)$ denote the result of applying f exactly y times to x
- For each i we store $f^{[2^j]}(i)$ as a sparse table

Binary lifting ctd.

- Let $f^{[y]}(x)$ denote the result of applying f exactly y times to x
- For each i we store $f^{[2^j]}(i)$ as a sparse table
- Then we can compute these in increasing order of j ,
calculating $j = 1$ using f itself and then for larger j letting
$$f^{[2^j]}(x) = f^{[2^{j-1}]}(f^{[2^{j-1}]}(x))$$

Binary lifting ctd.

- Let $f^{[y]}(x)$ denote the result of applying f exactly y times to x
- For each i we store $f^{[2^j]}(i)$ as a sparse table
- Then we can compute these in increasing order of j , calculating $j = 1$ using f itself and then for larger j letting $f^{[2^j]}(x) = f^{[2^{j-1}]}(f^{[2^{j-1}]}(x))$
- Thus we can precompute the table in $\mathcal{O}(n(\mathcal{O}(f) + \log(n)))$ and each query takes $\mathcal{O}(\log(m))$, a much better time complexity

Sparse table example

7	1	6	4	8	0	9	2	2	7	1	6
---	---	---	---	---	---	---	---	---	---	---	---

$j = 0$

Sparse table example

7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8											
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7										
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10									
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12								
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8							
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8	9						
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8	9	11					
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8	9	11	4				
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8	9	11	4	9			
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8	9	11	4	9	8		
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8	9	11	4	9	8	7	
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

8	7	10	12	8	9	11	4	9	8	7	6
7	1	6	4	8	0	9	2	2	7	1	6

$j = 1$

$j = 0$

Sparse table example

18	19	18	21	19	13	20	12	16	14	7	6
↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑
8	7	10	12	8	9	11	4	9	8	7	6
7	1	6	4	8	0	9	2	2	7	1	6

$j = 2$

$j = 1$

$j = 0$

Sparse table example

37	32	38	33	35	27	27	18	16	14	7	6
↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑
18	19	18	21	19	13	20	12	16	14	7	6
8	7	10	12	8	9	11	4	9	8	7	6
7	1	6	4	8	0	9	2	2	7	1	6

$j = 3$

$j = 2$

$j = 1$

$j = 0$

Sparse table example

$$\text{query}(1, 8) = 19 + 9 + 2$$

37	32	38	33	35	27	27	18	16	14	7	6	$j = 3$
18	19	18	21	19	13	20	12	16	14	7	6	$j = 2$
8	7	10	12	8	9	11	4	9	8	7	6	$j = 1$
7	1	6	4	8	0	9	2	2	7	1	6	$j = 0$

Sparse table example

$$\text{query}(0, 9) = 37 + 9$$

37	32	38	33	35	27	27	18	16	14	7	6	$j = 3$
18	19	18	21	19	13	20	12	16	14	7	6	$j = 2$
8	7	10	12	8	9	11	4	9	8	7	6	$j = 1$
7	1	6	4	8	0	9	2	2	7	1	6	$j = 0$

Example problem: Stikl

- <https://open.kattis.com/problems/stikl>